



BNB CHAIN L2

BEP-20 Native Token: OMC

# Omniverse Compute (OMC)

Decentralized GPU Computing Network — Whitepaper Deck v3.0

---

# Three ideas behind OMC

- Honest scope. Inference, single-node fine-tuning, batch rendering and other embarrassingly parallel jobs — not frontier-scale cluster training.
- Verification you can price. A per-job policy: spot-check re-execution, optional full redundancy, TEE where available, staking and slashing.
- Token demand tied to usage. Fees paid per job, part of every fee buys back and burns OMC, rewards are paid only against verified work.

| Layer                 | Role                                                             |
|-----------------------|------------------------------------------------------------------|
| L1 — BNB Chain        | OMC token, staking, DAO governance, final arbitration            |
| L2 — OP Stack (opBNB) | Heartbeats, job records, state-channel settlement, micropayments |
| Off-chain grid        | Sandboxed containers with CUDA / ROCm acceleration               |

## PROBLEM

# Compute is scarce, expensive and centralized

### Demand

AI developers face high hourly prices, regional quota limits and vendor lock-in; studios face sharp render peaks.

### Supply

Consumer GPUs, small data-center idle capacity and post-PoW hardware sit underused.

### The gap

No widely trusted way for strangers to rent GPU time with guarantees about correctness, privacy and payment.

### The answer

OMC: a DePIN market with verifiable jobs, USD-denominated staking and per-job settlement on BNB Chain.

## TARGET USERS

# Who runs what on OMC

### AI startups & developers

Inference, LoRA/QLoRA fine-tuning, batch embeddings — lower cost, no quota waits.

### Web3 & DeAI projects

Agent inference and batch workloads with crypto-native payment.

### Rendering & 3D studios

Blender and UE offline rendering with per-job pricing and burst capacity.

### GPU owners & small DCs

Monetize idle hardware with predictable earnings and low operating friction.

# Scheduling in two stages

- Stage 1 — hard filters: VRAM, driver and CUDA/ROCm version, region, confidential-computing requirement, minimum stake and reputation.
- Stage 2 — scoring: each metric becomes a percentile rank among eligible nodes, then a weighted matching problem maximizes total score; the final pick is drawn at random from the top-k candidates.

| Job class                      | TFLOPS | VRAM | Bandwidth | Latency | Reputation |
|--------------------------------|--------|------|-----------|---------|------------|
| Batch and rendering (proposed) | 0.30   | 0.20 | 0.15      | 0.10    | 0.25       |
| Inference (proposed)           | 0.20   | 0.15 | 0.10      | 0.30    | 0.25       |

# No single technique verifies every workload

| Mechanism                    | How it works                                                     | Cost to requester                         |
|------------------------------|------------------------------------------------------------------|-------------------------------------------|
| Spot-check re-execution      | A random share of jobs is silently re-run on an independent node | Base $\times$ (1 + a); proposed a = 5–10% |
| Full redundancy              | The job runs on 2–3 independent nodes; outputs are compared      | 2 $\times$ to 3 $\times$                  |
| Confidential computing (TEE) | Hardware-protected execution with remote attestation             | Premium for TEE nodes                     |
| Zero-knowledge proofs        | Research track — not offered at launch                           | —                                         |

Non-determinism is handled with task-specific rules: exact hashes, numerical tolerance, or fixed seeds.

## Honest about what cannot be guaranteed

- Isolation: hardened sandboxes (gVisor by default, Kata / microVM tiers for sensitive workloads), no host access, no outbound network by default.
- Encryption: inputs, outputs and images encrypted in transit and at rest; keys released only after assignment.
- Three data-protection levels: Standard, Protected (ephemeral, wiped), Confidential (TEE-only).
- Honest limitation: on non-TEE nodes a malicious provider could observe data in GPU memory — sensitive workloads should choose the Confidential level.
- Assurance: audits before mainnet with public reports; bug bounty from mainnet (critical findings from \$50,000, proposed).

# Fee flows tied to real usage

| Flow                        | Split                                                         |
|-----------------------------|---------------------------------------------------------------|
| Protocol fee at settlement  | 3% of every job                                               |
| Buy back and burn           | 30% of the fee = 0.9% of settled volume                       |
| Reputation Reward Pool      | 70% of the fee, paid for verified work                        |
| Payments                    | OMC (10% fee discount, proposed) or USDT / USDC on BNB Chain  |
| Emission scaling (proposed) | daily emission = ceiling × min(1, U / Utarget), Utarget = 60% |

0.9%

of settled volume is burned

3% fee × 30% burn share

\$10M

monthly volume → ~\$90,000 burn

illustration only, not a forecast

## TOKEN DISTRIBUTION

# 1,000,000,000 OMC fixed supply (BEP-20)

| Category                        | Share | Amount      | Vesting                                            |
|---------------------------------|-------|-------------|----------------------------------------------------|
| Compute mining and node rewards | 50%   | 500,000,000 | 0% at TGE; 10-year emission, halving every 2 years |
| Institutional investors         | 10%   | 100,000,000 | 5% at TGE, 6m cliff, 24m monthly                   |
| Private investors               | 8%    | 80,000,000  | 10% at TGE, 3m cliff, 18m monthly                  |
| Community airdrop               | 2%    | 20,000,000  | 100% at TGE                                        |
| Ecosystem and grants            | 15%   | 150,000,000 | 5% at TGE, 6m lock, 36m monthly                    |
| Team and core contributors      | 10%   | 100,000,000 | 0% at TGE, 12m cliff, 24m monthly                  |
| Treasury and liquidity          | 5%    | 50,000,000  | 50% at TGE for liquidity; rest by DAO Treasury     |

Circulating supply at TGE  $\approx$  65.5M OMC (6.55%). A public vesting dashboard shows every category.

## EMISSION

# A ceiling, not an entitlement

| Epoch | Years | Emission ceiling | Share of pool |
|-------|-------|------------------|---------------|
| 1     | 1–2   | 258.06M          | 51.6%         |
| 2     | 3–4   | 129.03M          | 25.8%         |
| 3     | 5–6   | 64.52M           | 12.9%         |
| 4     | 7–8   | 32.26M           | 6.5%          |
| 5     | 9–10  | 16.13M           | 3.2%          |

- Rewards are paid only against verified work.
- Unused emission goes to an Emission Reserve — burned or re-released only by DAO vote.

# Security priced in USD, penalties tiered

- Stake = Base\_Deposit (\$200, proposed) × tier multiplier ÷ OMC TWAP — set in USD so it does not drift with volatility.
- Tiers (proposed): consumer 1× · workstation 2× · data-center (80GB+) 4×.

| Offense                                | Penalty (% of stake, proposed) |
|----------------------------------------|--------------------------------|
| Job timeout                            | 2%                             |
| Unexpected offline during a job        | 5%                             |
| Incorrect output (failed verification) | 30%                            |
| Proven fraud or fake heartbeat         | 100%                           |

Slashed tokens: 50% compensate requesters · 30% DAO Treasury · 20% burned.

## From team-led to community-led

- Launch: 4-of-7 multisig of core contributors with a 3-day timelock (proposed); signers not disclosed at this stage.
- Transition: OMC holders vote on parameters and Treasury spending — proposed: 100,000 OMC threshold, 10% quorum, 7-day voting, 2-day delay.
- Full DAO (2027 or later): multisig removed or limited to emergency actions.
- Phase 1 (Q1–Q2 2026, completed): testnet on BNB Chain (Chapel), token contract, open-source client, 2% airdrop.
- Phase 2 (Q3 2026 – Q1 2027, in progress): mainnet targeted for Q1 2027, emission start, Hugging Face and Blender integrations, audits.
- Phase 3 (2027 Q2 and later): TEE tier, ZK research results, DAO transition, cross-chain settlement.

LEARN MORE

## Where to verify everything

### Website

<https://omc.network> — news, whitepaper v3.0, testnet, airdrop

### DCF forum

<https://forum.omc.network> — GPU pricing research and deep dives

### Repository

<https://github.com/CBTdao/omc-network> — open-source node client

### Contact

[support@omc.network](mailto:support@omc.network) · [security@omc.network](mailto:security@omc.network)

Third-party market figures are sourced and dated in Whitepaper v3.0, Section 13.